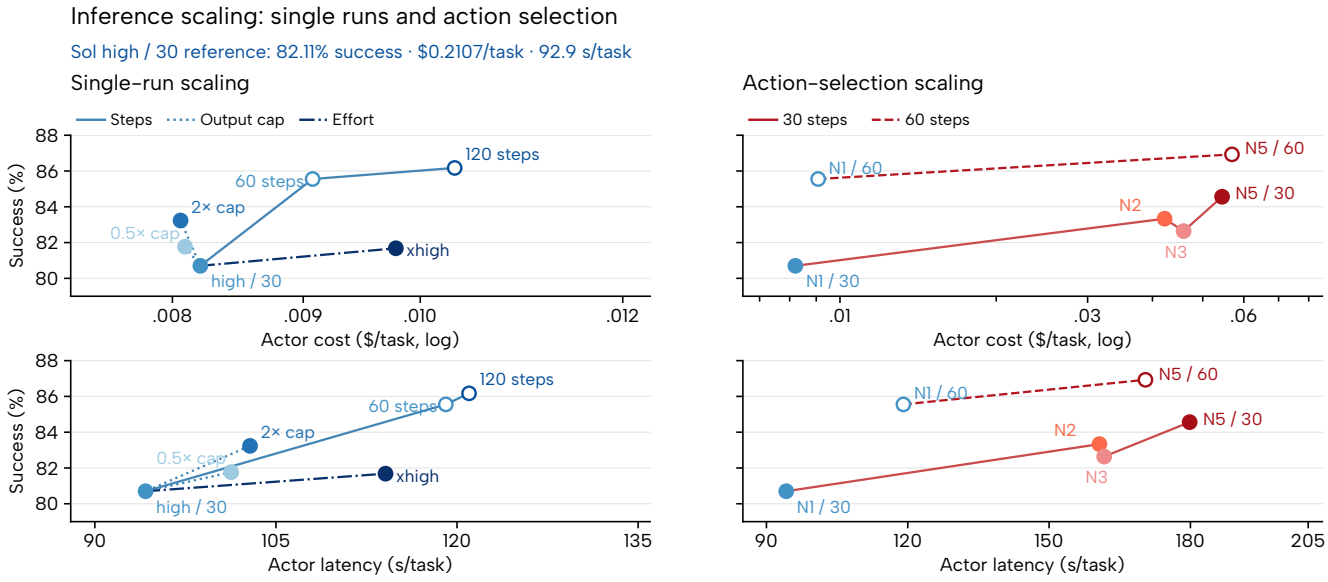
Inference scaling: single runs and action selection
Sol high / 30 reference: 82.11% success · $0.2107/task · 92.9 s/task
Single-run scaling
Steps
Output cap
Effort
Success (%)
88
86
84
82
80
120 steps
60 steps
2× cap
0.5× cap
high / 30
xhigh
.008
.009
.010
.012
Actor cost ($/task, log)
Action-selection scaling
30 steps
60 steps
N5 / 60
N1 / 60
N5 / 30
N2
N3
N1 / 30
.01
.03
.06
Actor cost ($/task, log)
Success (%)
88
86
84
82
80
120 steps
60 steps
2× cap
0.5× cap
high / 30
xhigh
90
105
120
135
Actor latency (s/task)
N5 / 60
N1 / 60
N5 / 30
N2
N3
N1 / 30
90
120
150
180
205
Actor latency (s/task)
10 Luna settings + Sol reference · Same 744 valid tasks · Terra judged · Planned benchmark weights
Local x-axis ranges · Single/60 advertised 30 · Live windows differ · Actor metrics exclude judging