Inference scaling: single runs and action selection

Sol high / 30 reference: 82.11% success · $0.2107/task · 92.9 s/task

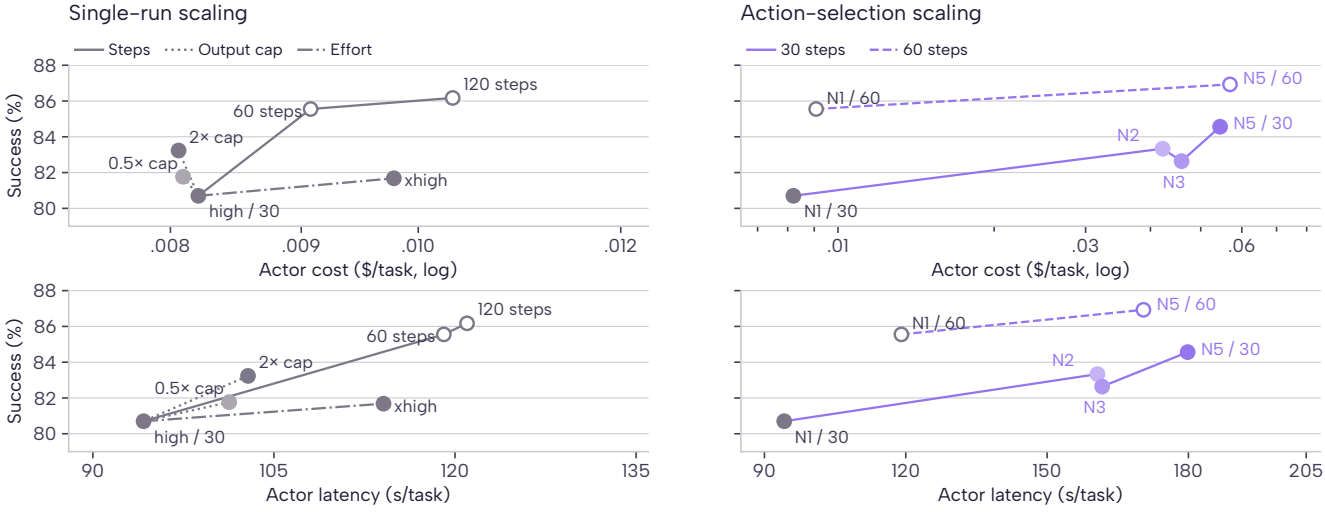

10 Luna settings + Sol reference · Same 744 valid tasks · Terra judged · Planned benchmark weights

Local x-axis ranges · Single/60 advertised 30 · Live windows differ · Actor metrics exclude judging