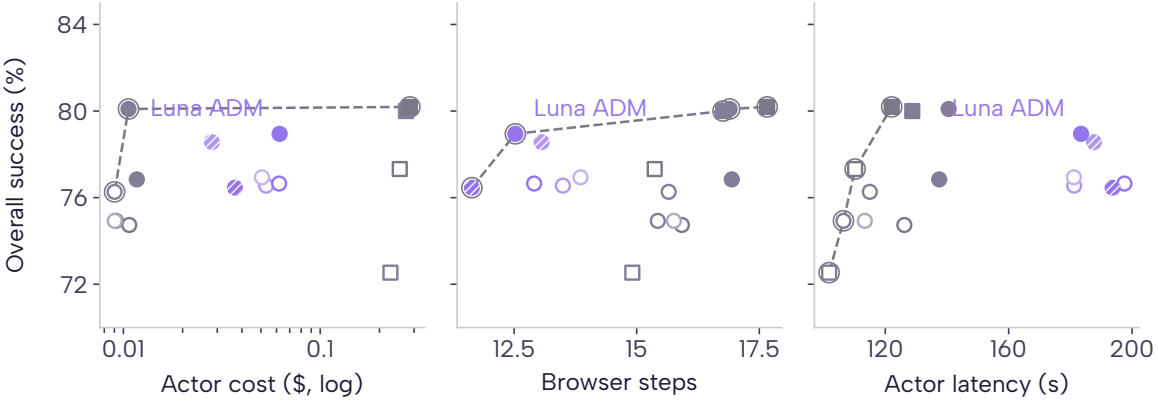

Single-call scaling (N=1)

- Luna baseline
- Luna xhigh effort
- 0.5× max tokens
- 2× max tokens
- Luna 60 steps
- Luna 120 steps
- Sol xhigh · 30 steps
- Sol high · 30 steps
- Sol high · 60 steps
- Sol high · 120 steps

Best-of-N action decision scaling

- N=2 · 30 steps
- N=3 · 30 steps
- N=5 · 30 steps
- Luna ADM N=5, 60 steps
- N=3/30 cached
- N=5/60 cached

Observed frontier · Cached context

Overall = successes / 1,045.
Invalid/missing count as nonsuccess.
X axes: valid-task actor means.
GPT-6 high unless xhigh. N=candidates.