Inference scaling after one-pass gap recovery

Sol high / 30 reference: 80.59% · $0.2197/task · 98.2 s/task

Single-run scaling

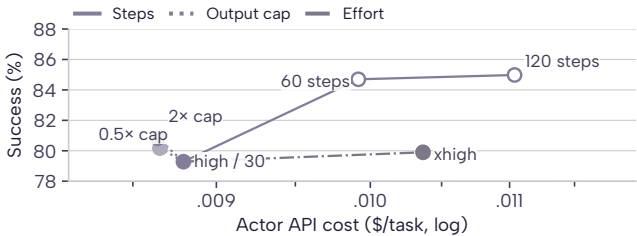

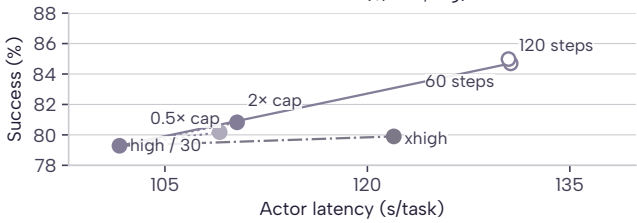

Action-selection scaling

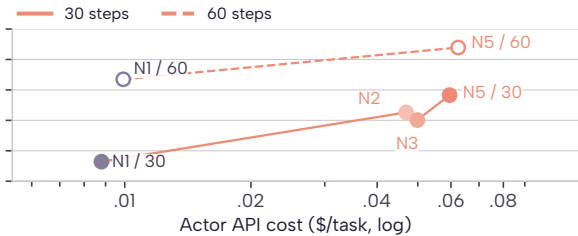

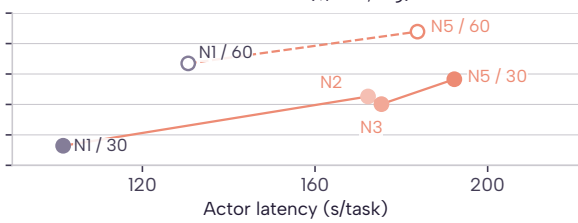

Same 953 common-valid tasks: OM2W 252 · WebVoyager 556 · DeepShop 145

Fixed planned weights 300/595/150 · Local x ranges · Actor metrics exclude Terra

Original valid outcomes retained · Single/60 advertised 30 · Live windows differ · Cached/Sol60 excluded