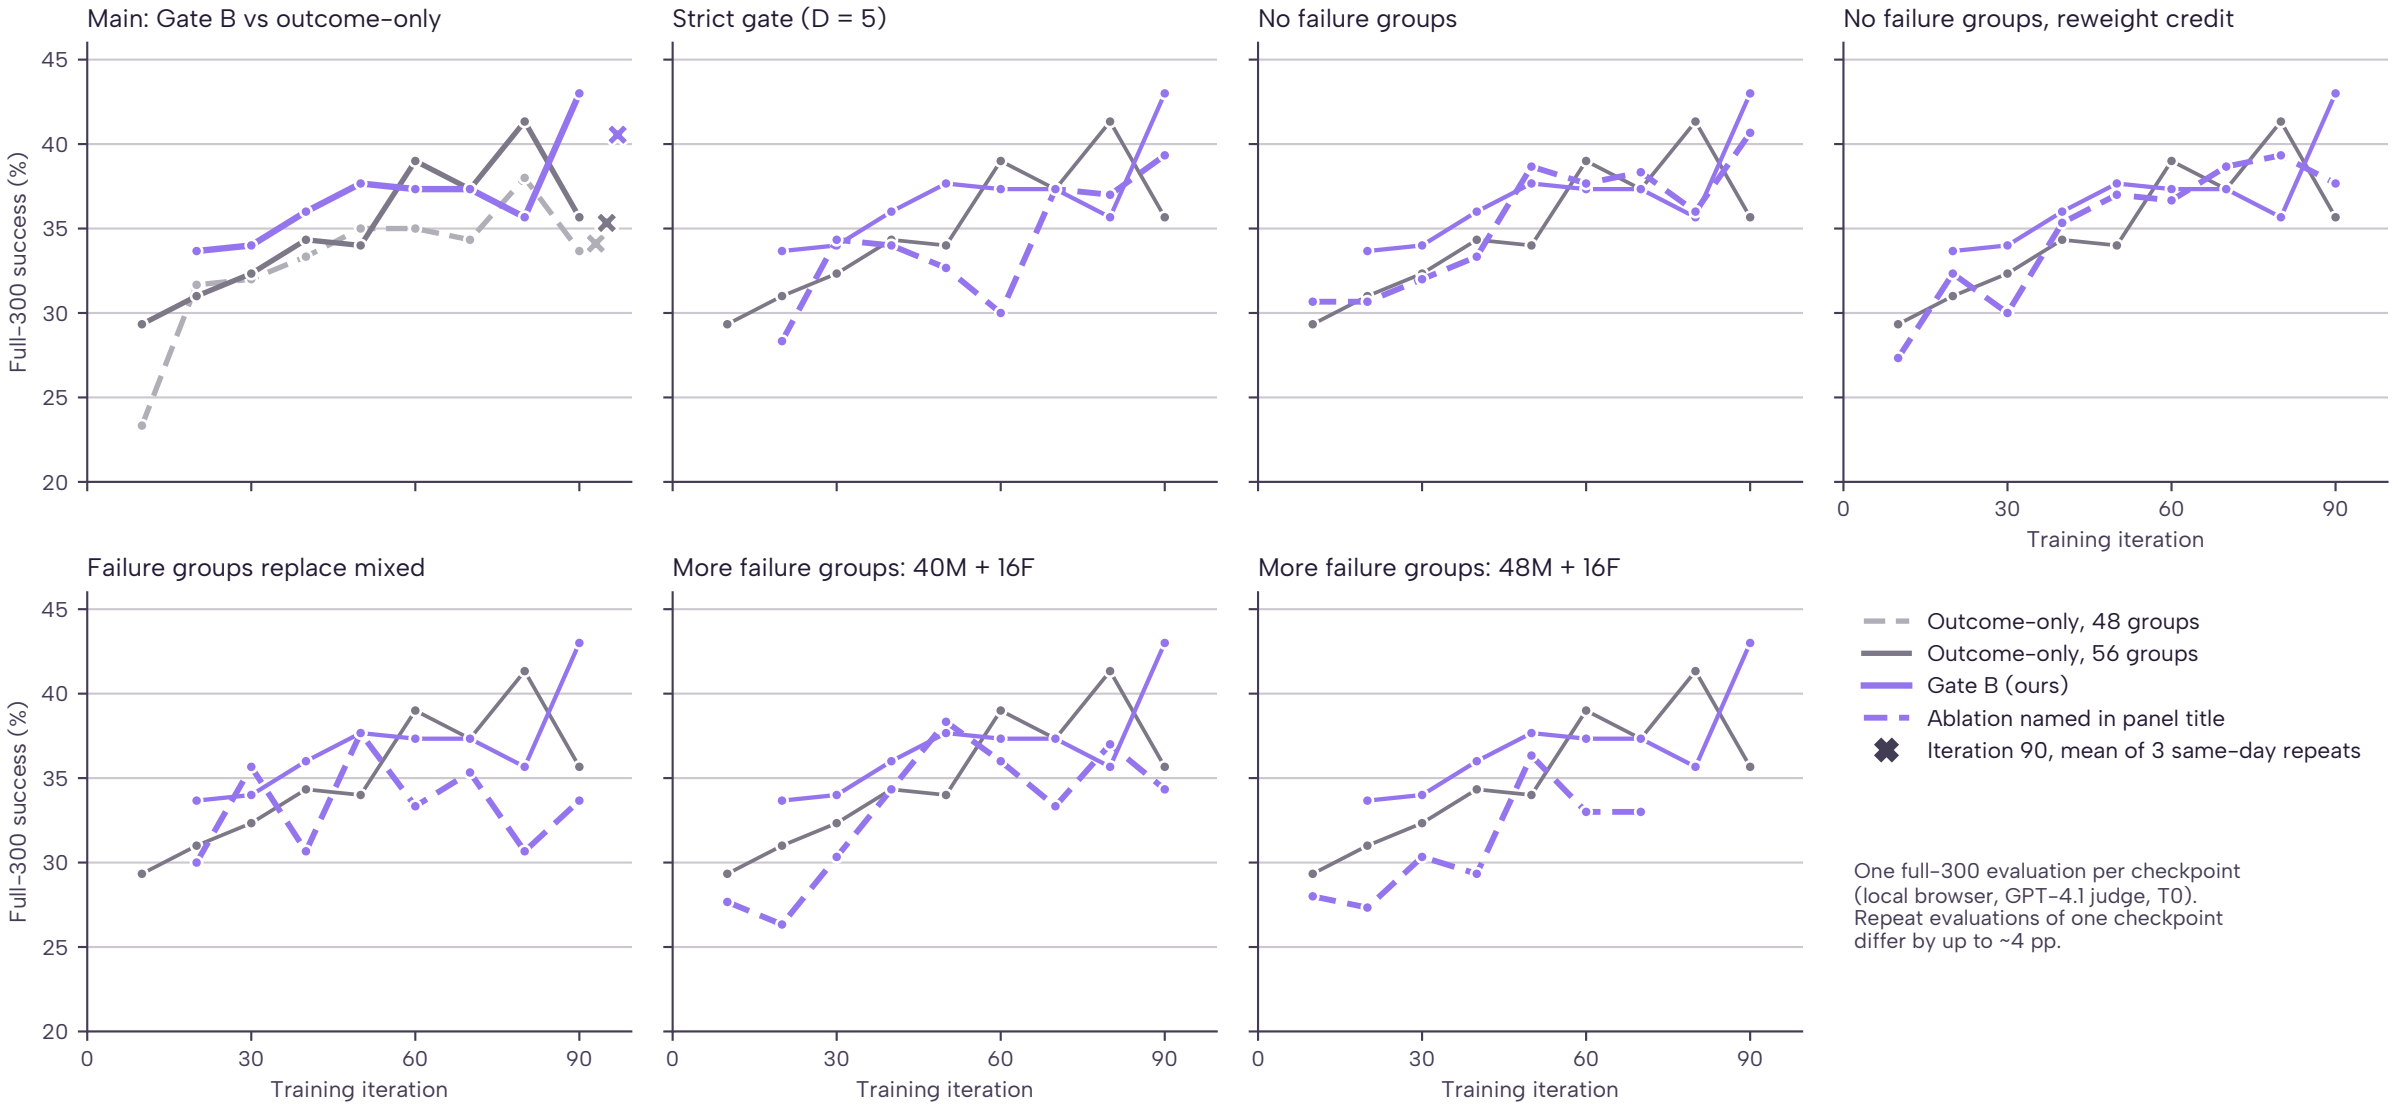
Main: Gate B vs outcome-only
Strict gate (D = 5)
No failure groups
No failure groups, reweight credit
Failure groups replace mixed
More failure groups: 40M + 16F
More failure groups: 48M + 16F
Full-300 success (%)
Training iteration
20
25
30
35
40
45
0
30
60
90
Outcome-only, 48 groups
Outcome-only, 56 groups
Gate B (ours)
Ablation named in panel title
Iteration 90, mean of 3 same-day repeats
One full-300 evaluation per checkpoint
(local browser, GPT-4.1 judge, T0).
Repeat evaluations of one checkpoint
differ by up to ~4 pp.