Selection heads trained on Piotr's panels: success per benchmark, labelled with the gain over one SFT 4B sample

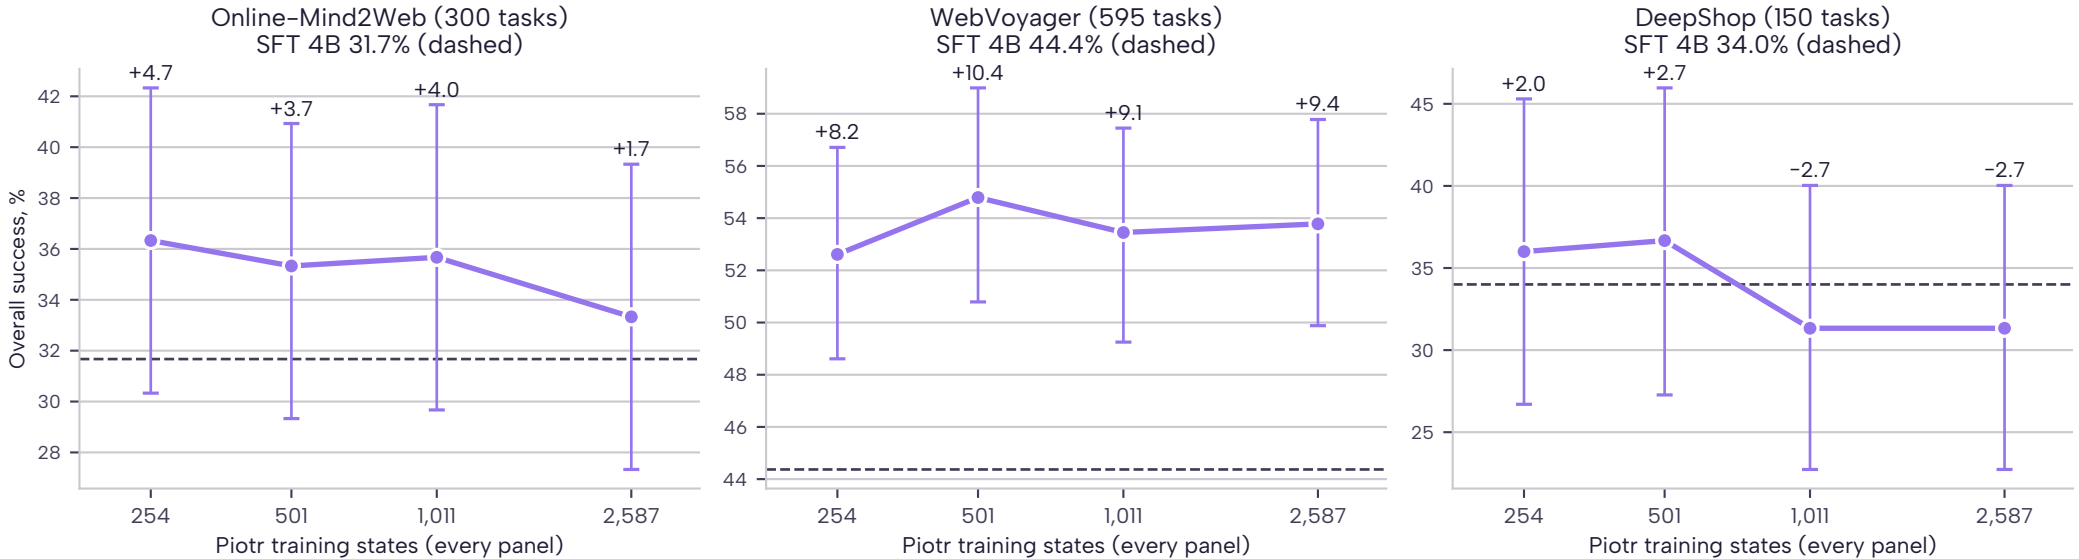

Bars: paired 95% bootstrap of the gain over SFT 4B on the same tasks, drawn around the head's rate; dashed: SFT 4B. One training seed per point; heads run 10-23/24, SFT 4B 10-18 (DeepShop drifts between dates). Online-Mind2Web judged by gpt-5.6-terra, the others by GPT-4o.