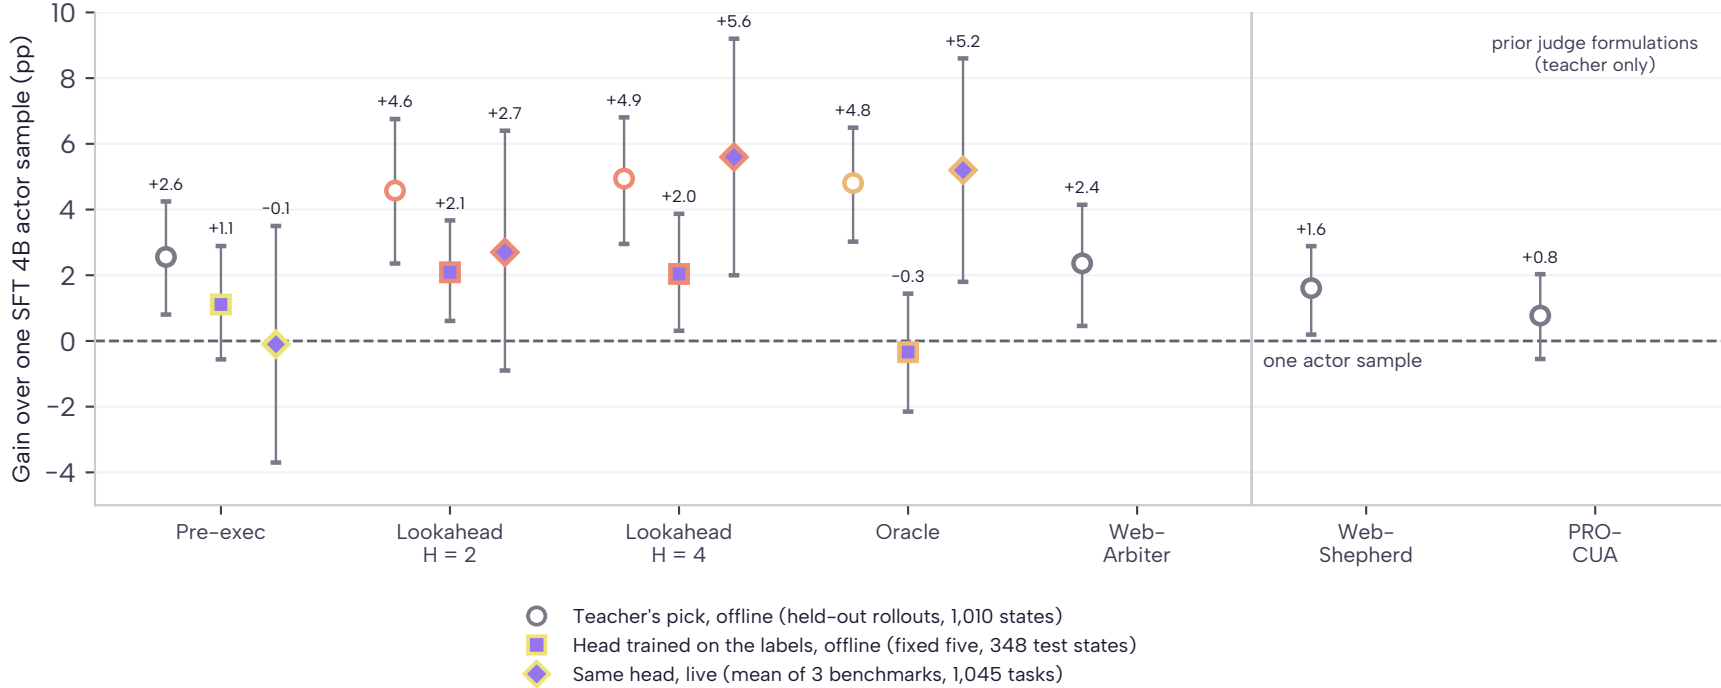

Bars: paired 95% intervals (bootstrap over tasks). Stages use different states and tasks: compare orderings within a stage, not magnitudes across stages.
One training seed per head. Oracle = best action on rollouts 1–3 (observed outcome); its head is trained on those labels. Rollouts numbered 1–8 as in the methods figure.