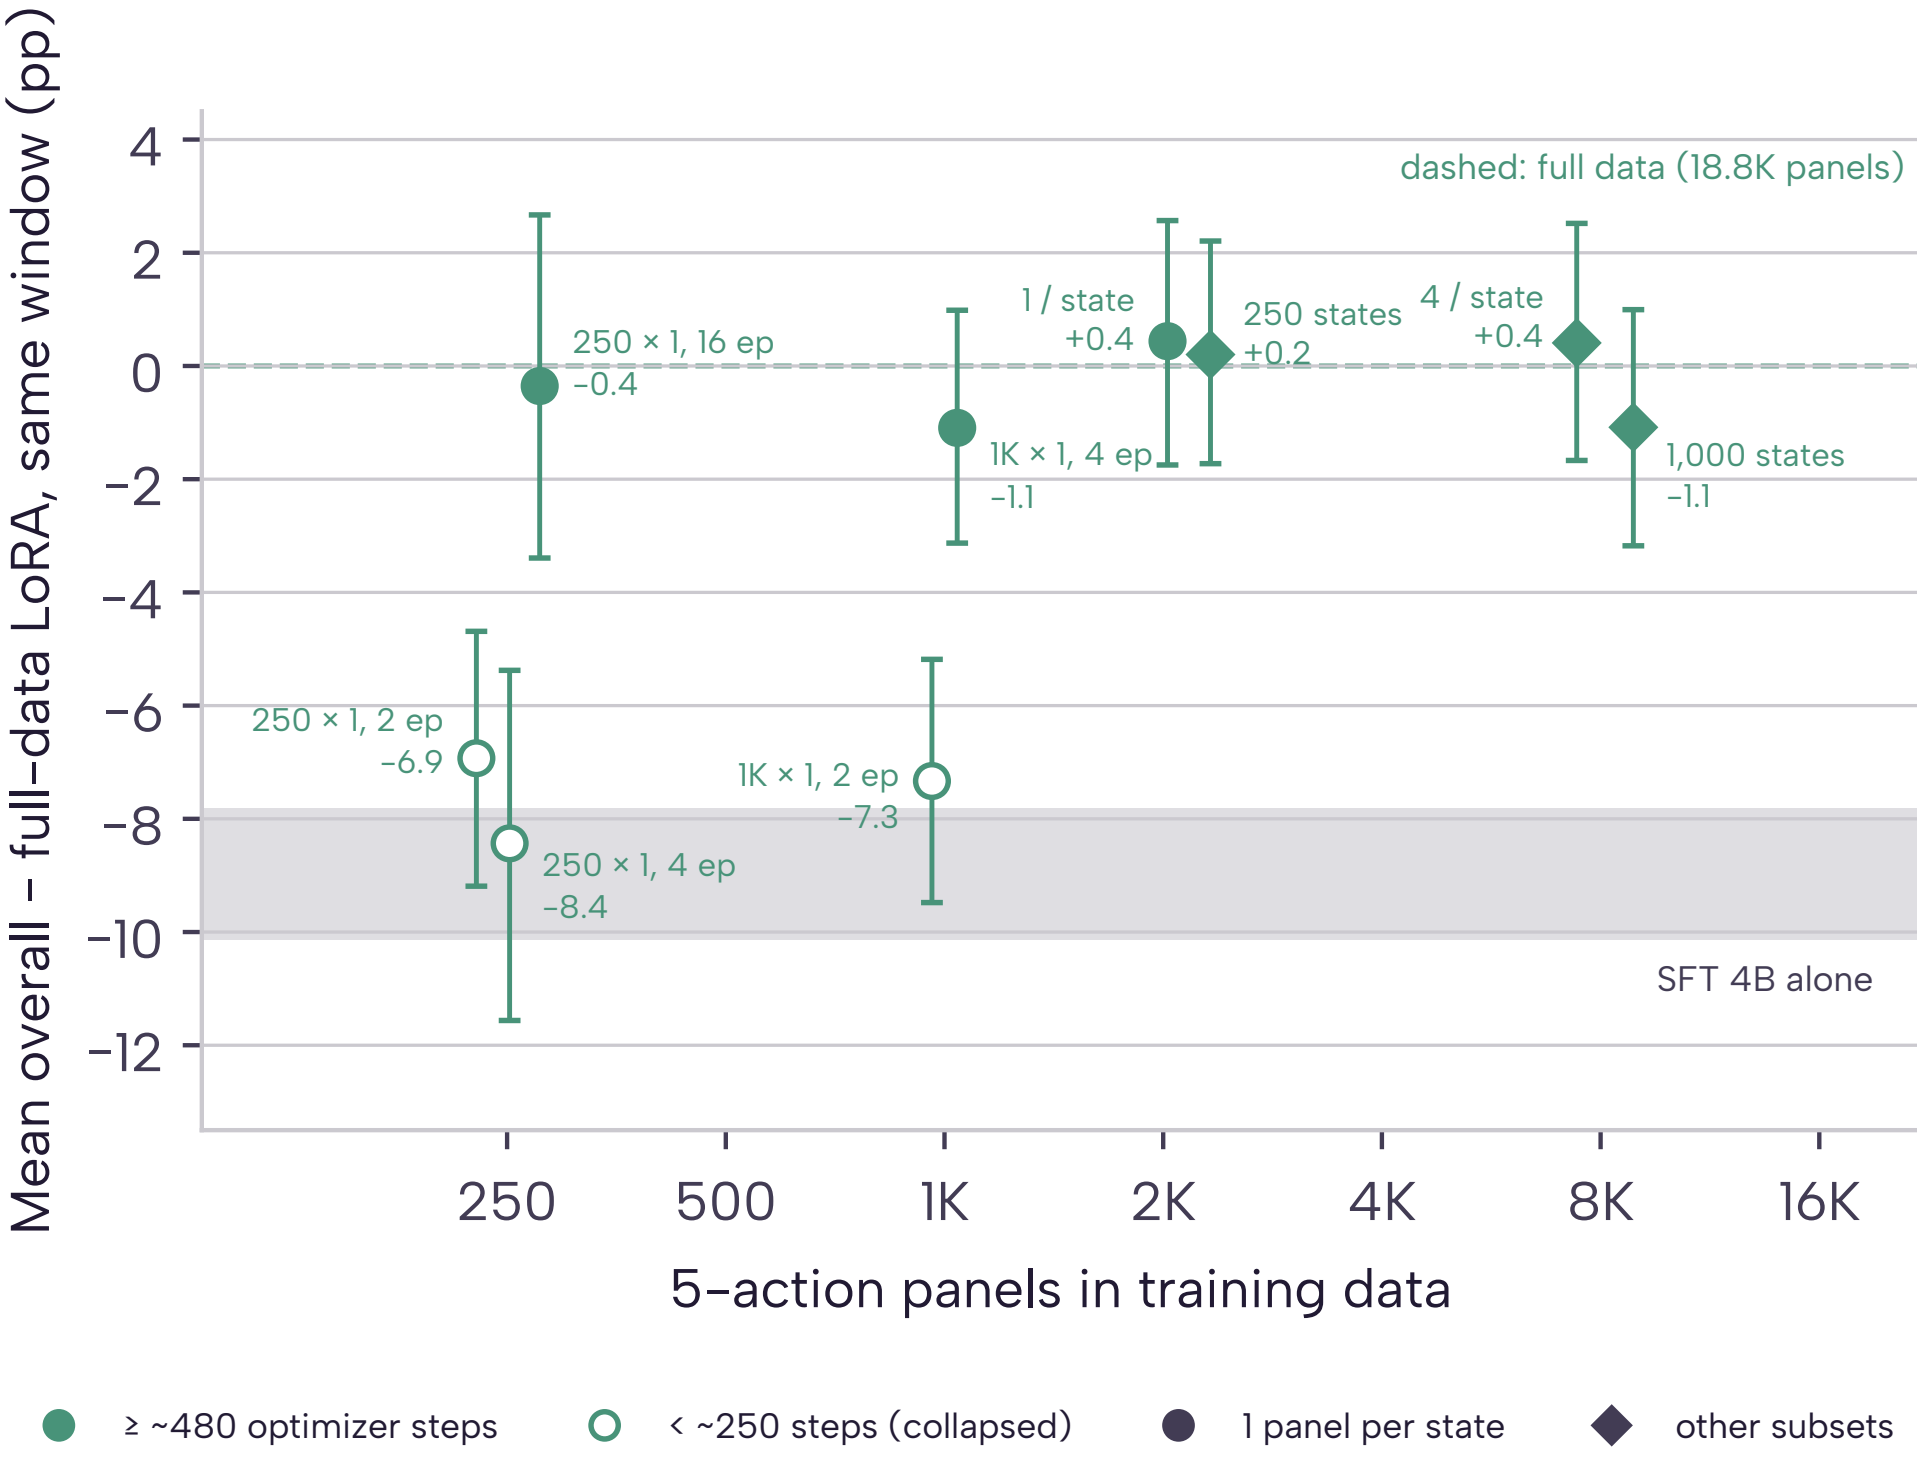
Mean overall − full-data LoRA, same window (pp)
4
2
0
−2
−4
−6
−8
−10
−12
dashed: full data (18.8K panels)
250 × 1, 16 ep
−0.4
1K × 1, 4 ep
−1.1
1 / state
+0.4
250 states
+0.2
4 / state
+0.4
1,000 states
−1.1
250 × 1, 2 ep
−6.9
250 × 1, 4 ep
−8.4
1K × 1, 2 ep
−7.3
SFT 4B alone
250
500
1K
2K
4K
8K
16K
5-action panels in training data
≥ ~480 optimizer steps
< ~250 steps (collapsed)
1 panel per state
other subsets