Selection heads trained on one GPT-5.5 panel per Piotr state: success per benchmark, labelled with the gain over one SFT 4B sample

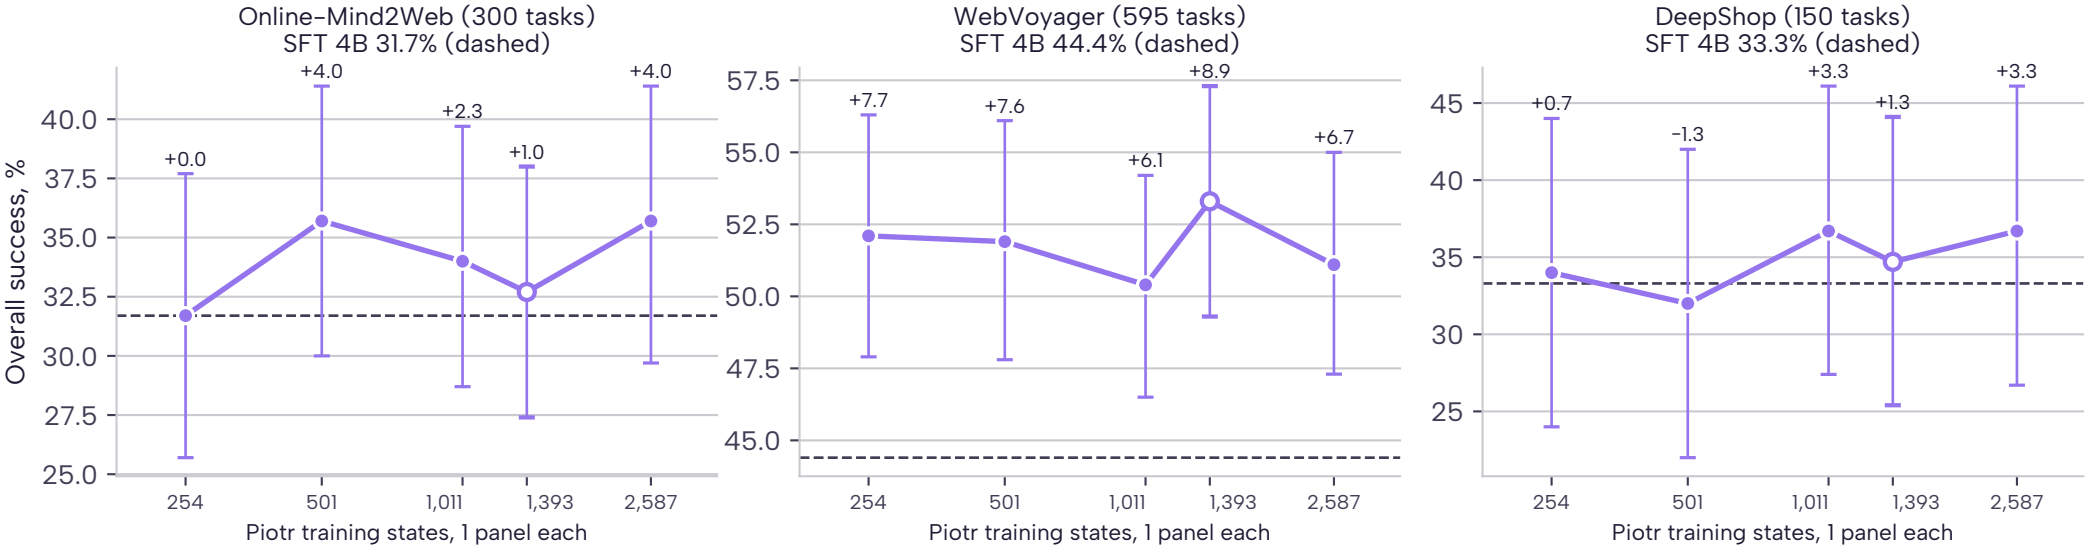

Bars: paired 95% bootstrap of the gain over SFT 4B on the same tasks, drawn around the head's rate. One training seed per point.
Hollow: earlier one-panel head (OM2W/WV run 10-23/24, DeepShop 10-26); others 10-26/27. DeepShop vs same-night SFT 4B (10-27).